


8/22

Test Cases: Validity, Originality, Comment

- Print error if fails

OSDev.org

Different languages are good at different things

in so long, "

- We can't build an OS in Python

c = getch();

- ↳ Reads next char and returns
 - What does reading next char mean?
 - Reads from standard input

↳ Choose file, keyboard, network connection to remote system, etc.

- You don't care about the source abstraction so you don't care abt file as input even tho completely different

Abstraction

- Hide details - simplicity
- Make things look the same - uniformity

Process - Running a program
Program - putting it in memory
Bytes on disk

Resource Management

- OS controls which programs get what memory, which program gets input, etc.

Isolation / Protection

- Program can't read another's files w/ correct permission

Inter-process Communication

- Processes share info w/ each other

Magic
Piece

UX - User Experience

- what non-programmers see - outside aspect - windowing, drag + drop, etc.

Connection b/w process and magic piece

`C = getch();`

- If user leaves w/out entering a char, what happens?
 - Has loop to check for char - wasting memory
 - Branch predictor becomes very efficient at running loop - fills pipeline w/ useless instructions, slows down / drains battery of system
- Want to **suspend** a program
 - Process can have multiple things happening - **Concurrency**
 - If waiting for input, other things can still be happening
 - Parallelism is a subset of concurrency
 - Concurrent things can run sequentially or parallel
 - Program can have concurrent things happen, system can have concurrent processes
- Concurrency - `async/await`
 - Async packages code into future and return, can evaluate/do wtv

3 abilities

- Suspend one thing
 - Maintain list of everything else to keep running
 - Resume suspended thing - detect conditions to resume
-

getch(callback)

- Callback function
 - When you get a character, run callback function
- Very simple concurrency
- Classic JS requires callback
- Hard to see sequential program - hard to follow, lots of jumping around
 - Modern JS/Python give illusion of sequential - asynchronous code splits into callbacks
- If need to pass variables would need to pass into callback parameter, callback needs to pass into other callbacks when itself

To share a variable, compiler needs to capture var and pass it around

- Need context of function - environment

⚠ Closure - Piece of code + environment - variables it needs

- Python - lambda, nested functions (functions inside functions)

int x = ...;

getch([x])(char ch){ } - Using closures

- [] w/ nothing to left means its a closure
- [x] means telling compiler to capture value of x
 - Shallow copy only of value of x
 - Can capture x by using &x - use pointer instead
 - Dangerous to memory at address x might get mapped to something else (x might not live long enough)
- Still looks ugly, ifs + loops suck, lots of nesting closures

Rule: Don't pass pointers or references unless you're convinced pointer will live long enough

Future<char> c = getchar(); - want non-blocking I/O survive what

- Instead of returning char, returns another data structure called a future
- Queues up functions so runs in correct order (if have 3
getch() functions)

Taking piece of work, package up into data structure, place into queue

- ↳ All ways of implementing getch() are doing this in essence

8/24

Don't have standard C++ library since running directly on hardware

- Std lib come from kernel from OS, we're creating OS

C++ is a combination of Java and C

- Method vs functions

- Methods bound to data, come from classes

- Functions are pure code

```
struct ABC {
```

- int a;

- float b;

- void m1() { }

```
}
```

struct vs class

- everything private by default in class

- everything public by default in struct

```
void func() {
```

- ABC x; - In C++, x is not a pointer - it's an actual instance of ABC

```
}
```

methods don't take space in a struct

```
void func() {
```

- ABC *x = new ABC(); - Pointer to ABC, creating space on heap for it

```
}
```

(Can use auto instead of type)

class ABC {

$ABC()$ ξ. ξ- Constructor

ABC (double z) §...§

`ABC x = {};` - initializing using default constructor \Rightarrow Not creating pointer

ABC p \approx 20.0% would be destroyed after function ends

auto p = new ABC(20.0) - Creating ptr, created in heap

i can use delete p;

new = null

`delete` == `free`

class ABC {

`~ABC()` - Destructor - Called when delete used object is destroyed

三

class Thing {

int vj

public:

Thing (int v) {

this. $v = v_j$

۳

class Thing {

int v_j

public

3

二

$$\text{Thing}(\text{int } x) = v(x) \in \mathbb{S}$$

ones a role:
initializer of v equals π

sane as
left side,
just more efficient
code

class Thing {

Overload + sign

int operator +(int z) { return v + z; }

}

Thing a { };

print(a + 20); - prints int

class Thing { overload function call, obj pretends to be function - returns int
int operator () (char c) { ... }
}

Thing a { };

a('a');

closure which captures v - will not run, stored to stack
int v=100;
auto a = [v] (char c) { print(c+v); };

need to copy to
heap

Later calls a('x')

↓

Compiler creates class which stores captured vars, overloads () operator
to store body of lambda/function call

Closure passed by value

getch([v] (char c){...}) - Closure inside copied into getch, original destroyed since only stored in stack

void getch()

takes closure, but we don't know the type of a closure

Use generic in Cpp w template

↳ function generator

template <typename T>

T silly(T a, T b){ return a + b; }

template <typename T>

void getch(T closure){ ... closure('x'); }

Thread - chunk of sequential code that runs concurrently w other threads

- share global variables and heap

- have own private variables (unshared)

- template, takes lambda as argument

thread([]{ ...});

1

Mental model:

- Thread() function returns immediately, continues to run sequential code
- Code inside lambda also runs, run concurrently
 - Can be in parallel, whenever - not sure

How to implement thread?

Queue of things to run

8/29

KernelInit called by one core, wakes up the rest

Hypervisor flag - Detect if running on real hardware or VM

thread [E] { } - Run concurrently w/ whoever calls it
- Both running at the same time

Kernel - Piece of software - first thing to run on system
- Most fundamental piece of OS

- Only 1 per system - Manages threads

Core - Physical piece of hardware to execute instructions

Thread - Software abstraction, virtual - cores create threads to run concurrent code

One core can call thread and another can run the work inside thread

void kernelMain(void) {

1. printf("Start")

Within one thread, 1 → 2 → 3

2. thread [] {

- Inside the inner thread, I → II

I. printf("Mid")

- I will happen before II

II. printf("Mid2") }

- However, we don't know if III will run before or after I / II

3. printf("End")

- II will be called before I (thread must be created first)

Yield() - Thread says on the thread can take over core

- Scheduling point - Can come back to run on another core

Hyperthreading - core can run multiple hardware threads

- In our world, one core can only run one thread at a time

step() - give someone else a chance to run

- Same as yield but doesn't come back

void step() {

 while (true) yield(); - Very bad implementation, keeps running, doesn't
 } free memory

 - Functions correctly, just wasteful of resources

void thread (Work work) {

 work();

 - Makes same core run work(), removes
 parallelism

 - Every core needs to do work

 L Create queue to store threads to run

 - Ready Queue - Threads ready to run but not running

 - What's the best way to schedule threads? Priority?

Queue has thing to go into queue, type of lock for queue

Better to define global variables in .h file using extern, defined in .cc file

- Preprocessor substitutes induces textually - if defined in h, multiple definitions

- Static limits access to only one file

Control Block - Block of memory to control things

- Called TCB for threads

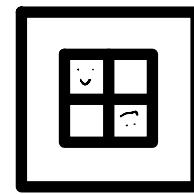
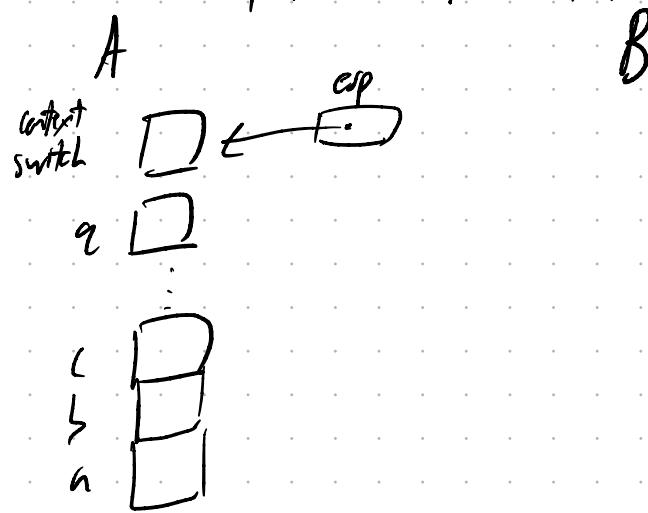
Think hard about using address - data can get destroyed
↳ Use $TCB \oplus Tcb = \text{new } TCB$;

Link elements in queue together using next pointer in TCB

8/3

Need to preserve context in thread switching

- Maintain pointer to top of stack



Threads share heap and global variables

- Only local variables are private

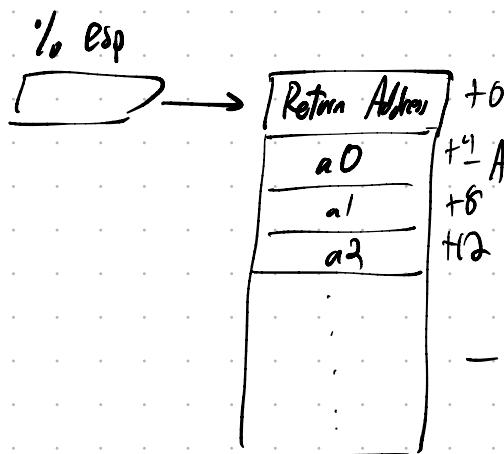
- Save registers as well

Caller-saved - Function called can use them

1. eax, edx, ecx, ebx, esi, edi, ebp, esp ↗ Stack pointer
Must be preserved, are callee-saved

8 registers, ignoring float/double/vector registers

2. args



+4 AV 32 bits a0...n are arguments in stack

- Caller pushes arguments into stack in reverse order

Can't have stack pointer saved into stack to try to preserve

- Would get updated to point to itself

- Instead, have context-switch() function take in pointer as argument

- This is where to store the stack pointer

Every thread has its own stack

context-switch(uint32_t* save-esp-here, uint32_t* restore-the-esp-here):

// (cc) dst
mov 4(%esp), %eax

- Save stack pointer address (dst arg) to eax reg

mov 8(%esp), %ecx

- Save and arg

// Using A's stack

push %ebx

push %esi

push %edi

- Save stack pointer into address given (address is in eax register)

push %ebp

mov %esp, (%eax)

// Saved my context

mov (%ecx), %esp

pop %ebp

pop %edi

// Using B's Stack

pop %esi

pop %ebx

ret

Expect stack to have 5 items: popped registers and return address
Stack grows toward low memory addresses

Return address is pointer to next

$$100 = (01, 9)$$

8/7

Lack of **atomicity** - Two cores can increment a variable at the same time, both set to the same value

Functions that start w -- - **intrinsic function**

- Defined in compiler

- **atomic-fetch-add()**

- Fetch (load) value and add to it

- Do atomically so no one can come in between

- Processors run things out of order - $A \rightarrow B \rightarrow C$ can be run as $A \Rightarrow C \rightarrow B$

- If B is atomic, has chance C will run first and load in value before B incremented due to speculation

- Will observe changes in value in wrong order

- **ATOMIC-SEQ-CST** - Forces sequential ordering

- **Weak memory ordering** - If one core runs $X \rightarrow Y$,

- other cores may not see it in order, can see new X and old Y be loaded out of order

write() ≈

```
ready = true;
```

```
val = 10;
```

```
}
```

```
read()
```

```
while (!ready);
```

```
print(val)
```

```
}
```

→ Will not work bc race condition → print val before set

write() {

```
    val = 10;  
    ready = true;  
}  
  
read() {  
    while (!ready);  
    print(val)  
}
```

Still might not work bc assumes strong ordering
- Compiler might speculate ready → true
- reads old value of val before changes

Enforce a **fence** - p1 cannot see changes after fence before changes
before fence happen

- Slows down system ten of cycles
- **Barrier** - close, used for coordinating threads, synchronization
- Fence for correct memory ordering

load_fence() - Make sure all loads before fence happen before loads after

- Only controls loads, bringing them memory

store_fence() - Make sure stores before fence happen before stores after

Usually need both

- Only writings to memory

ATOMIC SEQ-CST - Do whatever you can to ensure sequential ordering

```
write() {  
    some_value = 10;  
    store_fence();  
    value_is_ready = true;  
    //value_is_ready = true;  
}  
  
reader() {  
    while (!value_is_ready);  
    load_fence();  
    print(some_value);  
}
```

- add-fetch - add, then return new value
 - fetch-add - get old value then add
-

Spinlock

Spinlock {

Atomic couts taken;

SpinLock() : take(fake) {}

SpinLock(const SpinLock &) = delete;

void lock(void) {

while (!take(exchange(true))) {

}

}

void unlock(void) {

take.set(fake);

}

}

Need to know if I own the lock - If I set fake → true or alt was true

↓ Exchange - set to value, return old value

Use atom

x = 10;

l.unlock(); - Guarantees sequential

y = 20;

LockGuard exists to protect against forgetting to unlock, leave what unlocking - destructor to help delete

Issue: No fairness, many threads racing to get lock could cause **starvation** - one thread could never get the lock

Bigger Issue: Lock() function delaying core its running on waiting resources, other threads can run when I'm waiting

Issue: Single core and two threads

Spinlock lock §3

Thread [1] §

lock.lock();

yield();

lock.unlock();

};

Thread [2] §

lock.lock();

};

1st thread locks

1st thread yields to 2nd thread

2nd thread tries to take the lock

- Sits there spinning indefinitely b/c
lock never released

lock() also wasted by branch speculation fill buffer, then needs to flush when actually set to false

Solutions

yield() inside while loop in lock()

- lets others run, solves deadlock problem

- Problem: yield() uses a Spinlock (ready queue has a lock)

- Becomes deadlock/runs out of memory

power() instruction in x86

- Says slow down, help save energy - Doesn't solve 2 thread problem

halt() instruction in x86

- Stronger power(), puts core in low-power state, stops fetching instructions until interrupt happens (timer, keyboard press, etc.)

- Way to save energy, not giving others a chance to run either

- Doesn't solve 2 thread problem

mwait0() instruction

- Give memory address of something don't like

- Wait until its value changes, puts in low-power state until value changes

To solve 2 threads problem - need another lock that has yield()

Race Condition

```
void lock(void) {
```

```
    while (!taken.exchange(true)) {
```

```
        mwait(&taken);
```

```
}
```

```
{
```

If another thread unlocks here, mwait is waiting for the wrong value

Solution

```
monitor() instruction monitor(&taken)
```

- Say plan on waiting for memory location, tell hardware
- When we mwait(), registered interest w/ hardware - can detect before start waiting - if changes before hits mwait(), change next mwait() to MPl.

```
void lock(void) {
```

```
    monitor(&taken);
```

```
    while (!taken.exchange(true)) {
```

```
        mwait();
```

```
        monitor(&taken);
```

```
}
```

```
{
```

Have another lock that uses yield()

```
void lock(void) {  
    while (!taken.exchange(true)) {  
        yield();  
    }  
}
```

yield gives chance for anyone holding the lock to release

Problems:

- Unfair, could have starvation
- Inefficient - still asking question repeatedly, just wasting time in between
 - Still waiting time

Solution

Add "blocked" status to ready and running states

Blocking Lock - Stop running completely until lock is unlocked

- New queue for blocked threads, don't run until unlocked

- Attach queue to lock

- Have flag and queue as variables in lock

- Allows us to stop running + checking for unlock

- Solves unfairness thru queue w/ our policy for popping/running threads

Create yield w/ queue to add to ~called Block(addr)

void unlock (void) {

auto who-needs-to-run = waiting.remove();

if (who-needs-to-run == nullptr)

taken.set(false);

else

ready-giveaway.add(who-needs-to-run);

}

Tons of race conditions - what to do?

- Add Spinlock inside Blocking lock

void lock(void) {

inner-spin-lock.lock();

while (!taken.exchange(true)) {

inner-spin-lock.unlock();

block();

inner-spin-lock.lock();

}

- Still has race condition as before

↳ If add after block() - no one to unlock since context switch

{}

L Pos: Spinlock as argument to block(), after switch + safely invoke
block-giveaway, then unlock

void unlock (void) {

inner-spin-lock (locked);

auto who-needs-to-run = waiting.remove();

if (who-needs-to-run == nullptr)

taken.set (false),

else

ready-queue.add (who-needs-to-run),

inner-spin-lock.unlocked;

}

waiting-q 0

|

taken true