

---

---

---

---

---

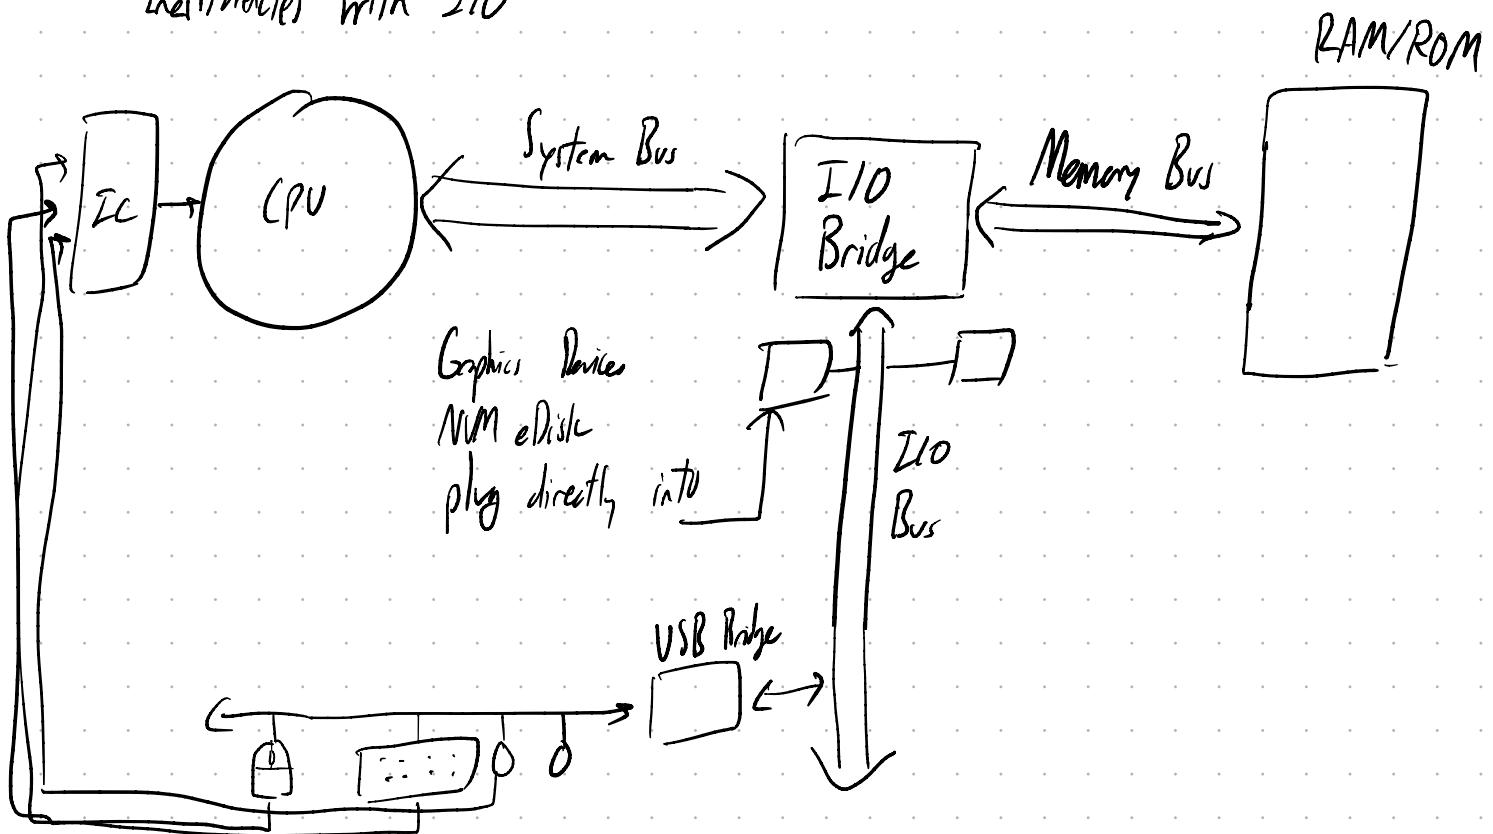


9/12

## Little Book of Semaphores

Processor Fetch Decode Execute Update PC

- Loop - don't point to next instr, point back to some place
- Inefficiencies with I/O



USB Bridge speaks cheaper language

If you click key on keyboard, what happens?

- USB Bus → I/O Bus → System Bus
- 1st Opt

- Can pull keyboard and ask if input

- Expensive, most likely doesn't have - humans are slow

- Wasting CPU time, bus, etc

- 2nd

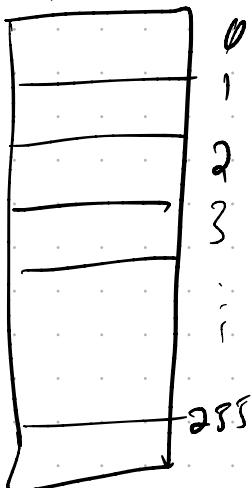
- Logical wire directly from keyboard to CPU - Still kinda expensive
- If has something change state of wire
- Sending message through bus
- Wire just has info keyboard was instructed then core goes and asks keyboard for key/keys
- Logical wire just says there was interrupt which device interrupted
- Can't have as many wires to core as number of devices - variable # of bus
- Instead one device can decide what device generated interrupt based on its number into core - **Interrupt Controller**

fetch decode execute - Sends interrupt to core are doing info about device number

FDX cycle stops normal cycle, fetches from custom place when interrupted

- All done by hardware

256 types of interrupt **IDT** - interrupt descriptor table



ids of interrupt types

- Pointer to piece of code that knows how to handle each type of interrupt - **ISR** - interrupt service routine

IC sends id to core, core follows ISR to fetch info  
**IDTR** (pointer to IDT) saved in memory

- Points to beginning and saves size  
ISR pointer also carries mode to run in (16-bit, 32-bit, etc.)

ISR is a function in our case (return context, current mode, flags)

ZFS typically pushes everything onto stack to diff know would be interrupted

Otherwise conflict b/w registers in previous process and ISR

- No risk of overflowing stack if interrupt during interrupt handle

- Usually just disable interrupt during interrupt handle

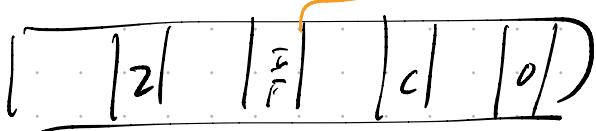
---

By time running ISR, interrupt disabled until return

- How to disable then enable?

x86

EFLAGS (N2CV in ARM) interrupt flag



- 0 Disabled for core
- 1 Enabled for core

pushf pushes all flags onto stack

popf pops

(Can use in context-switch - just need to change stack alignment)

Who to enable/disable interrupt?

tcb = active[SMP::mem()

- Can be interrupted during call

- Get SMP::mem() value, then interrupted before indexing into array, then yield, then TCR can get picked up from another core, then (SMP::mem) value is wrong

SpinLock - Interrupt interaction

- Both protect against race conditions

Interrupts can bring race conditions even to single-core - Spinlock cannot protect

→ edr x9,  
      add x9,x9, #1 — Interrupted here, interrupt handle modifies same X9, X9 might  
      str x9

↓  
Use interrupt safe lock

↳ Disable interrupt around critical conditions

Disable interrupt protects core against itself, Spinlock protects against other cores

(Both need to be O(1) otherwise becomes unresponsive)

↳ Sometimes will need both

disable()

spin\_latch()

Not O(1)

↳ Spinning can take arbitrarily long, can disable for  $\infty$  long

[ Critical O(1) ]

spin\_munlock()

enable()

spin\_latch()

disable()

Can get interrupted in between

- Can get deadlocked

[ Critical O(1) ]

enable()

spin\_munlock()

ISL - Specialized lock to combine

- While spinning, disables then enables interrupt

- 1) Don't use locks if possible
- 2) Use Blastinglock() if possible - can't use in yield() or inside Blastinglock
- 3) Use ISL

9/14

Implement TSL to solve problem w/ Spinlock + Interrupt disables

```
    disable();
    lock
    while (flag.exchange(true)) { }   Atomics swap to true
    enable();
    disable; — looks like cancels out, but enable allows pending interrupts
}
}
```

Unlock flag.set(false)

thread[i] { ... }

stuff here is work done on behalf of the thread

↓  
imagine calls yield() → void yield() { ... }

At some point, code is for core to do work for itself

- when messng w/ ready queue, not thread that's ready but core is running

Interrupts are similar

- When thread runs, pit runs, causes interrupt, pitHandler() runs
- Thread stops running, doing own housekeeping stuff, doing core work

Who can call yield()? The thread or the core?

- Core is always calling everything b/c
- Threads are yielding, not cores
- Yield means let another thread run, core itself doesn't yield

stop();

(can add core thread in init)

- Thread thing, cores don't stop

thread is just 0x

Spinlocks - uses

while(thread){}

ESL - Use over Spinlocks

delete stuff();  
yield();

BlockingLocks - Don't spin, instead has queue to add self to

33

- Threads add self to queue when blocked

- In block, add yourself to waiting queue and context switch to another thread
- Cores can't block, threads block

init.c → correct place - check - no thread → thread shows up out of nowhere

- calls stop, malloc which uses BlockingLock

threads using Spinlocks are wasteful

- Cores can spin, if many and nothing else to do, can just spin

threads shouldn't disable interrupts - should be able to be preempted

Only threads can block, keep uses BlockingLock, only threads can do sleep things

Interrupt Handler runs w/ interrupt disabled

- Doing a core thing, or calling yield() is funky
- Behaving like a thread in context where doing core-only things
- Why problem?

- Shouldn't be trying to get Spinlock when interrupts disabled

- Calling yield() from the heap, which uses lock

- Cannot do lots of work w/ interrupts disabled

Apit Handler (interrupt handler) has interrupt disabled

- Call yield in the middle - yield is not short duration
- Not good

## Solutions

Don't do complex work in interrupt handler

- Off-level interrupt handle: Real interrupt handler runs really fast, when return, returns to another handler w/ interrupt handler w/ interrupts enabled to do more complex stuff instead of where it was
- In theory can get interrupted in off-level handler + never return back to where it was

Don't use TSL in ready queue

- Use ready queue per core, don't need TSL()
    - Can have imbalance where one core can have many things to do, another has none
    - loses parallelism
  - Background activity to balance queues
- 

## Blocking lock

Have busy flag, if true, don't spin but add self to waiting queue

- Then switch to another thread + give chance to own
- Details
  - What if try to add to queue then switch + ready queue is empty?
  - Know why come back - bc lock is unlocked or ready queue is empty?
  - Can have idle thread - make initial TCB idle threads

- what if X gets lock + succeed, A tries to get lock, is not taken  
add to queue, X releases lock after A tries to lock + before A adds  
self to queue - X sees queue is empty, doesn't see A is there
- flare ESL for this
  - Care is spinning, not thread

Have nothing queue waiting for all other threads to be added - Blocking barrier